

The Interpreter Project

CSC 413 - Software Development

Anthony J Souza

Sourced from Dr. Levine

Anthony J Souza ©2023

Computer Science
San Francisco State University

Contents

1	Introduction	3
2	The Interpreter	3
2.1	Frames (Activation Record) and the Runtime Stack	3
3	Supported ByteCodes	6
3.1	Halt ByteCode	6
3.1.1	Requirements	6
3.1.2	Dumping	6
3.2	Pop ByteCode	7
3.2.1	Requirements	7
3.2.2	Dumping	7
3.3	FalseBranch ByteCode	8
3.3.1	Requirements	8
3.3.2	Dumping	8
3.4	Goto ByteCode	9
3.4.1	Requirement	9
3.4.2	Dumping	9
3.5	Store ByteCode	10
3.5.1	Requirements	10
3.5.2	Dumping	10
3.6	Load ByteCode	11
3.6.1	Requirements	11
3.6.2	Dumping	11
3.7	Lit ByteCode	12
3.7.1	Requirements	12
3.7.2	Dumping	12
3.8	Args ByteCode	13
3.8.1	Requirements	13
3.8.2	Dumping	13
3.9	Call ByteCode	14
3.9.1	Requiriements	14
3.9.2	Dumping	14
3.10	Return ByteCode	15
3.10.1	Requirements	15
3.10.2	Dumping	15
3.11	Bop ByteCode	16
3.11.1	Requirements	16
3.11.2	Dumping	16
3.12	Read ByteCode	17
3.12.1	Requirements	17

3.12.2 Dumping	17
3.13 Write ByteCode	18
3.13.1 Requirements	18
3.13.2 Dumping	18
3.14 Label ByteCode	19
3.14.1 Requirements	19
3.14.2 Dumping	19
3.15 Dump ByteCode	20
3.15.1 Requirements	20
4 Interpreter Class	21
5 ByteClassLoader Class	22
5.1 ByteClassLoader Functions	22
6 Program Class	23
6.1 Program Class Functions	23
7 RuntimeStack Class	25
7.1 RunTimeStack Class Functions	26
8 VirtualMachine Class	28
9 Interpreter Dumping	29
9.1 DUMPING IMPLEMENTATION NOTES	31
10 Coding Hints	32
10.1 Possible Order of Implementation	32
10.2 Other Coding Hints	32
11 Setup	34
11.1 Importing Project	34
12 Requirements	35
13 Submission	36

1 Introduction

For this assignment you will be implementing an interpreter for the mock language X. You can think of the mock language X as a simplified version of Java. The interpreter is responsible for processing byte codes that are created from the source code files with the extension x. The interpreter and the Virtual Machine (this will be implemented by you as well) will work together to run a program written in the Language X. The two sample programs are a recursive version of computing the nth Fibonacci number and finding the factorial of a number. These files have the extension x.cod.

And as always, if there are any questions please ask them in discord AND in class. This promotes collaborative thinking which is important. NOTE THIS IS AN INDIVIDUAL ASSIGNMENT but you can collaborate with other students.

2 The Interpreter

In the following sections you will find coding hints, class requirements, and other information to help you implement and complete the interpreter project. Please read these pages a few times before you begin coding.

2.1 Frames (Activation Record) and the Runtime Stack

Frames or activation records are simply the set of variables (actual arguments, local variables, and temporary storage) for each function called during the execution of a program. Given that function calls and returns happen in a LIFO fashion, we will use a Stack to hold and maintain our frames.

For example, if we have the following execution sequence:

<pre>Main Call A from Main Call B From A</pre>
--

With the sequence above, we will get the following runtime stack(3):




Figure 1: View of runtime stack as methods are called.

Now if we take this concept and apply it to some code written in the Language X we can see a better example of our runtime stack.

```

program { int i int j
    int f ( int i ) { 2
        int j int k 3
        return i + j + k + 2 4
    }
int m 1
m = f(3) 5
i = write(j+m)
}

```

Figure 2: Small code sample in the Language X.




Figure 3: View of Runtime Stack as code is executed.

In the figure above we can see what happens as a program in the X language is executed. Below is a general idea of what is going on in the virtual machine.

- At Bubble 1 - When code reaches this point, the virtual machine has pushed 3 values, `i`, `i` and `k`. All three have the values of 0.
- At Bubble 2 - When we reach this bubble, we have done a few things. First we need to load the literal value of 3 on the runtime stack. This will act as an argument for the function `f`. Create a new activation frame for function `f`, which includes the top value of the runtime stack since it takes one argument. Then we can call function `f`.
- At bubble 3 - After executing up to bubble 3, the virtual machine has pushed 2 more values to the runtime stack. This is to declare and initialize the variables `j` and `k`. NOTE: `j` in `main`'s activation frame and `j` in `f`'s activation frame DO NOT conflict.
- At Bubble 4 - After executing up to bubble 4, the virtual machine has evaluated the expression `i + j + k + 2` and pushed the result onto the top of the stack.
- At Bubble 5 - Once the virtual machine gets to bubble 5, it has done a few things. First, it saves the value 5 (the return value) and then clears the rest of function `f`'s activation record. Then removes the frame boundary for function `f`. Then pushes the return value back to the runtime stack. Then moves the value 5, into the slot on the runtime stack that stores the variable `m`.

3 Supported ByteCodes

Below will you will find the descriptions of the ByteCodes your interpreter must support. Each ByteCode must be implemented to the correct specifications outline below. It is important that your understand what each ByteCode does as this will help you implement the ByteCodes.

3.1 Halt ByteCode

The HALT ByteCode is used to alert the virtual machine that program execution is to be stopped. Halt is not allowed to kill the interpreter program. Therefore, Halt may not call `System.out.exit` to stop the execution of the program.

3.1.1 Requirements

- Request that the VirtualMachine halt execution of current program.
- Halt takes no arguments.
- Halt ByteCodes are not be Dumped.
- Halt cannot execute a `system.exit` function call.

3.1.2 Dumping

The Halt bytecode is to not be dumped.

3.2 Pop ByteCode

The Pop ByteCode will be used to remove values from the run time stack. The Pop ByteCode is not allowed to remove values across frame boundaries. It is the implementers responsibility to ensure that Pop is only allowed to pop the appropriate number of values. Pop has one argument, N , which is the number of values to be popped. You cannot make any assumptions about the value of N other than it is a strictly positive number.

3.2.1 Requirements

- Pop takes one argument which is the number of values to remove from the run time stack.
- Pop is not allowed operate across frame boundaries.
- If dump is on, Pop is required to be dumped.
- The Pop bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.2.2 Dumping

The Pop bytecode has the following dump syntax:

POP N

where N is the number of times pop is requesting to be executed. Note, this value can be different than the actual number of POPs executed depending on current frame size.

3.3 FalseBranch ByteCode

The FalseBranch ByteCode will be used to execute conditional jumps(think of executing control structures like if-statements and loops). FalseBranch will have one argument. This argument is a Label that will mark a place in the program to jump to. FalseBranch will remove the top value of the run time stack and check to see if the value is 0. If the value is 0, jump the corresponding label. If the value is something else, move to the next ByteCode in the program. FalseBranch will need to have its label address calculated before the program begins executing. This requires finding where the destination of the jump is going to be numerically(address in the program) in the program.

3.3.1 Requirements

- FalseBranch takes one argument, a label to jump to.
- FalseBranch's label address will need to be resolved. This requires computing where FalseBranch will jump to if the value popped from the stack is 0. Address resolution needs to be done before you began executing the program. This will be discussed later in this document.
- Remove the top value of the stack.
 - if value is 0, jump to label.
 - if value is not 0, move to next ByteCode. (do nothing)
- If dump is on, FalseBranch ByteCode is required to be dumped.
- The FalseBranch bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.3.2 Dumping

The FalseBranch bytecode has the following dump syntax:

```
FALSEBRANCH label
```

where label is the original label read from the .cod files.

3.4 Goto ByteCode

The Goto ByteCode is used to jump to Labels in our programs. Goto is considered an unconditional jump. This means regardless of the state of the program, we take the jump. Goto has one argument, the label it needs to jump to. Like FalseBranch, Goto's label needs to go through address resolution as well.

3.4.1 Requirement

- Goto has one argument, a label to jump to.
- Goto's Label must have its address resolved before the program begins executing. More on this later.
- If dump is on, Goto is required to be dumped. Examples are given in this document.
- The Goto bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.4.2 Dumping

The Goto bytecode has the following dump syntax:

```
GOTO label
```

where label is the originallabel read from the .cod files.

3.5 Store ByteCode

The Store ByteCode will be used to move values from the top of the run time stack to an offset in the current frame. This offset starts from the beginning of the frame. The idea behind this ByteCode is it is needed to do operations like assignments. The Store ByteCode is not allowed to operate across frame boundaries.

3.5.1 Requirements

- The Store ByteCode can have 1 to 2 arguments.
 - one argument is the offset in the current frame where the value that is popped is to be stored.
 - The second argument, if present, is the identifier (variable) the value being moved belongs to. This we be used for dumping.
- Store must pop the top of the runtime stack and store the value at the offset in the current frame.
- Store cannot operate across frame boundaries.
- Store's offset is strictly positive
- If dump is on, Store is required to be dumped.
- The Store bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.5.2 Dumping

The Store bytecode has the following dump format:

```
Basic Syntax : STORE <offset> [<id> <id>=<value-popped>]
<offset> is the index in the current from to load from
<id> is a variable identifier [optional]
Example      :
If we assume the RuntimeStack has the following values
BEFORE executing the Store bytecode:
[0,1,2,3]
Then executing a Store 1 k would produce:
STORE 1 k    k=3
```

note that if id is missing or not supplied than your dump should only print

```
STORE 1
```

3.6 Load ByteCode

The Load ByteCode will be used to move values from an offset in the current frame to the top of the stack. This offset works from the beginning of the frame. The purpose behind this ByteCode is it is needed to setup copies of values for things like expressions and arguments for functions. The load ByteCode is not allowed to operate across frame boundaries.

3.6.1 Requirements

- The Load ByteCode can have 1 to 2 arguments.
 - one argument is the offset in the current frame where the value is to be copied from.
 - The second argument, if present, is the identifier (variable) the value belongs to. This we be used for dumping.
- Load must copy the value at the offset in the current and push it to the top of the stack.
- Load must not remove any values from the runtime stack.
- Load cannot operate across frame boundaries.
- Loads's offset is strictly positive.
- Is it not required that the identifier argument exists.
- If dump is on, Load is required to be dumped.
- The Load bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.6.2 Dumping

```
Basic Syntax : LOAD <offset> [<id> <load id>]  
<offset> is the index in the current from to load from  
<id> is a variable identifier  
Example      : LOAD 2 j      <load j>
```

If id is missing then then dumping load will just be:

```
LOAD offset
```

therefore, using the above example, we have:

```
LOAD 2
```

3.7 Lit ByteCode

The Lit ByteCode is used to push literal values to the runtime stack. In some cases, Lit ByteCodes will be accompanied with an id (a variable name), this id represents the variable name the value belongs to. This id is optional.

3.7.1 Requirements

- The Lit ByteCode takes 1 or 2 arguments.
- The Lit ByteCode must push a value to the RunTimeStack.
- Is it not required that the identifier argument exists.
- If dumping is on, Lit ByteCode is required to be dumped.
- The Lit bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.7.2 Dumping

Basic Syntax : LIT <value> [<id> int <id>]
<value> is the value being pushed to RunTimeStack
<id> is a variable identifier [optional]

If we were to execute a LIT 0 j then the dump will be

```
LIT 0 j    int j
```

if the id is missing then lit has the following dump format:

```
LIT 0
```

3.8 Args ByteCode

The Args ByteCode is going to be used to setup how many arguments a function has. The Args ByteCode will always be executed just before a Call ByteCode. The Args ByteCode has one argument, the number of values that are arguments for the next function call. This value N, will be used to determine how many values starting from the top of the runtime stack will be a part of a newly created activation frame for the next function call. Args will need to figure out where in the runtime stack this new frame begins at and push this index into the FramePointer stack.

3.8.1 Requirements

- The Args bytecode has one argument, the number of values that will be a part of the new activation frame.
- The Args ByteCode will need to push the starting index of the new frame to the framePointer stack.
- If dump is on, the Args ByteCode is required to be dumped.
- The Args bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.8.2 Dumping

The Args bytecode has the following dump syntax:

ARGS n

where n is the number of args of the new call frame.

3.9 Call ByteCode

The Call ByteCode is what the VirtualMachine uses to jump to locations in the program to execute sections of code we call Functions. When encountered the Call ByteCode will jump to the corresponding label in the program. The ByteCde is also responsible of keeping track of where control should return to when a function completes its execution.

3.9.1 Requiriements

- Call ByteCode takes 1 argument, a label to jump to.
- Call Code must go through address resolution to figure out where it needs to jump to in the Program before the program is ran.
- Call Code must store a return address onto the Return Address Stack.
- Call Code must Jump the address in the program that corresponds to a label code (this address is computed during address resolution).
- If dumping is on, the Call ByteCode is requiried to be dumped.
- The Call bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.9.2 Dumping

The Call bytecode has the following dump syntax:

```
Basic Syntax : CALL <id>    <base-id>(<args>)
<id> is a function identifier
<base-id> is the actual id of the function;
e.g. for CALL f<<2>>, the <base-id> is f
<args> are the function arguments.
Example      :
If we assume the RuntimeStack has the following values
    [0,1,2] [3,4,5]
And we execute a CALL f<<3>>.
Before the CALL code is executed, an ARGS 3 has been executed.
Then the dumping of the call code looks as follows:
CALL f<<3>>    f(3,4,5)
```

3.10 Return ByteCode

The Return ByteCode will be used to return from functions but also to put return values in the correct position on the runtime stack. The interpreter project will use this convention for handling arguments and return values. Callers of functions are required to setup argument for the functions they call. Functions themselves (callees) are required to put return values in the correct spot just before returning from a function. Note that this is a convention we will adhere to and is something that is not enforced programmatically. This means if you fail to follow this convention, transient bugs can happen. The Return ByteCode has a lot of responsibility. The steps for completing a return is important.

3.10.1 Requirements

- The Return ByteCode can take 0 to 1 arguments. The arguments have no effect on its functionality. But does effect the Dumping process.
- The Return ByteCode must store the return value at the top of the runtime stack.
- The Return ByteCode must empty the current frame of all values when the function is complete.
- The Return ByteCode must pop the top value from the framePointer stack to remove the frame boundary.
- The return ByteCode must pop the top of the return address stack and save it into program counter.
- If dumping is on, the Return ByteCode is required to be dumped.
- The Return bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.10.2 Dumping

The Return bytecode has the following dump syntax:

```
Basic Syntax : RETURN <id>      EXIT <base-id>:<value>
<id> is a function identifier
<value> is the value being returned from the function
<base-id> is the actual id of the function;
e.g. for RETURN f<<2>>, the <base-id> is f
Example      : RETURN f<<2>>  EXIT f : 4
```


3.11 Bop ByteCode

The Bop ByteCode is used to implement binary operations for the Interpreter Project. The Bop ByteCode will need to remove 2 values from the runtime stack and operate on them accordingly. The result needs to be pushed back to the top of the stack. Be careful, the order of the operands matter.

3.11.1 Requirements

- Bop must pop 2 values from the runtime stack.
- Bop must push 1 value, the result, back to the top of the runtime stack.
- Bop must implement the following binary operations:

+-----+-----+			
Addition	: +	Subtraction	: -
Division	: /	Multiplication	: *
Equality	: ==	Not-Equal	: !=
Less-Than Equal	: <=	Less-Than	: <
Greater-Than Equal	: >=	Greater-Than	: >
Logical And	: &	Logical-OR	:
+-----+-----+			

- If dump is on, the Bop ByteCode is required to be dumped.
- The Bop bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.11.2 Dumping

The Pop bytecode has the following dump syntax:

BOP operator

where operator is the operation executed. For example, if we execute a + operation the, dumping the bytecode would be:

BOP +

3.12 Read ByteCode

The Read ByteCode is used to read user input from the keyboard. Only integers should be accepted from users. You may use Scanners or BufferedReaders to read input from the user.

3.12.1 Requirements

- When asking for user input, use the following prompt: "Please enter an integer : "
- The Read ByteCode needs to verify that the value given is actually a number. If an invalid number is given, state that the input is invalid and ask for another value. Continue to do so until a valid value is given.
- Push the validated integer to the VirtualMachine's RunTimeStack.
- If dumping is on, the Read bytecode is required to be dumped.
- The Read bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.12.2 Dumping

The Read bytecode has the following dump syntax:

READ

3.13 Write ByteCode

The Write ByteCode is used to display information to the console. The only thing Write is allowed to display is the top value of the runtime stack. No other information is allowed to be shown.

3.13.1 Requirements

- Prints the top of the runtime stack to the console.
- NO OTHER information can be printed by the Write ByteCode when printing the value.
- If dumping is on, bytecode is required to be dumped.
- The Write bytecode cannot detect when it should be dumped nor should it call dump in the VirtualMachine.

3.13.2 Dumping

The Write bytecode has the following dump syntax:

WRITE

3.14 Label ByteCode

The Label bytecode is a bytecode that has no functionality. Its sole purpose is to mark locations in the program where other bytecodes can jump to. Label bytecodes will be used in address resolution so other ByteCodes know where they are supposed to jump to.

3.14.1 Requirements

- Label takes one argument, a label which is used to denote a location in the program.
- If dumping is on, dumping the Label bytecode is optional.

3.14.2 Dumping

If you decide to dump the Label bytecode, it should have the following dump format:

```
LABEL label
```

where label is the string value that is the label of where to jump to.

3.15 Dump ByteCode

The Dump ByteCode is used to turn dumping ON and OFF. Dumping in the interpreter project is only done when dumping is ON.

3.15.1 Requirements

- The Dump ByteCode has 1 argument. Either "ON" or "OFF"
- The Dump ByteCode must request the VirtualMachine to turn dumping either "ON" or "OFF".
- The Dump ByteCode is NOT TO BE DUMPED.

4 Interpreter Class

The interpreter class is used as the entry point for this project. It is responsible for the following:

1. Reading in the .cod file to be ran from the command line.
3. Calling loadCodes to load the program.
4. Calling resolve address for the built program from previous step.
5. Creating the VirtualMachine and begin executing the program.

You can view the source code of the Interpreter class by looking in the Interpreter.java file in the project repository.

When running the program make sure to let your IDE/Text Editor know what program it is running by supplying a command line argument in the configuration settings.

5 ByteClassLoader Class

The ByteClassLoader class is responsible for loading bytecodes from the source code file into a data-structure that stores the entire program. We will use an ArrayList to store our bytexodes. This ArrayList will be contained inside of a Program object. Adding and Getting ByteCodes will go through the Program class.

The ByteClassLoader class will implement a loadCodes function that does the following:

1. Opens a file reader to parse the source code.
2. For each line in the source code, does the following:
 1. tokenizes the current line read
 2. With the first token on each line, instantiate a ByteCode instance.
 3. If there are more tokens, move them to an ArrayList of strings and pass these arguments to the bytexode's init function.
 4. Store the fully initialized ByteCode instance into the program data-structure.
3. returns program object.

5.1 ByteClassLoader Functions

```
1 /**
2  * Constructs ByteClassLoader object given a COD source code
3  * file name
4  * @param fileName name of .cod File to load.
5  */
6 public ByteClassLoader(String fileName){
7     this.codSourceFileName = fileName;
8 }
9
10 /**
11  * Loads a program from a .cod file.
12  * @return a constructed Program Object.
13  * @throws InvalidProgramException exception thrown when
14  * loadCodes fails.
15  */
16 public Program loadCodes() throws InvalidProgramException {
17     return null;
18 }
```

6 Program Class

The program class will be responsible for storing all the ByteCodes read from the source file. We will store ByteCodes in an ArrayList which has a designated type of ByteCode. This will ensure only ByteCodes and its subclass can only be added to the ArrayList.

6.1 Program Class Functions

```
1 /**
2  * Instantiates a program object using an
3  * ArrayList
4  * */
5 public Program()
6
7 /**
8  * Gets the size of the current program.
9  * @return size of program
10 */
11 public int getSize()
12
13 /**
14  * Grabs an instance of a bytecode at an index.
15  * @param programCounter index of bytecode to get.
16  * @return a bytecode.
17  */
18 public ByteCode getCode(int programCounter)
19
20 /**
21  * Adds a bytecode instance to the Program List.
22  * @param c bytecode to be added
23  */
24 public void addByteCode(ByteCode c)
25
26 /**
27  * Makes multiple passes through the program ArrayList
28  * resolving address for Goto, Call and FalseBranch bytecodes.
29  * These bytecodes can only jump to Label codes that have
30  * a matching label value. HINT: make note what type of
31  * data-structure ByteCodes are stored in.
32  * ~~~~ METHOD SIGNATURE CANNOT BE CHANGED ~~~~
33  */
34 public void resolveAddress()
```


Resolving symbolic addresses can be done in many ways. It is up to you to design a clean implementation for mapping the generated labels the compiler uses to absolute addresses in the program. An example is given below.

The Program class will hold the ByteCode program loaded from the file. It will also resolve symbolic addresses in the program. For example, if we have the following program below

```
0. FALSEBRANCH continue<<6>>
1. LIT 2
2. LIT 2
3. BOP ==
4. FALSEBRANCH continue<<9>>
5. LIT 1
6. ARGS 1
7. CALL Write
8. STORE 0 i
9. LABEL continue<<9>>
10. LABEL continue<<6>>
```

After address resolution has been completed the source code should look like the following (NOTE you should not modify the original source code file, these changes are made to the Program object):

```
0. FALSEBRANCH 10
1. LIT 2
2. LIT 2
3. BOP ==
4. FALSEBRANCH 9
5. LIT 1
6. ARGS 1
7. CALL Write
8. STORE 0 i
9. LABEL continue<<9>>
10. LABEL continue<<6>>
```

7 RuntimeStack Class

Records and processes the stack of active frames. This class will contain two data structures used to help the VirtualMachine execute the program. It is **VERY** important that you do not return any of these data structures. Which means there should be NO getters and setters for these data structures. These data-structures must remain private as well.

The RuntimeStack class will use the following two data-structures:

```
1 // This stack is used to record the beginning of
2 // each activation record(frame) when calling functions.
3 private Stack<Integer> framePointer
4
5 // This ArrayList is used to represent the runtime stack.
6 // It will be an ArrayList because we will need to access
7 // ALL locations of the runtime stack.
8 private ArrayList<Integer> runTimeStack
```

When trying to understand the Runtime stack class you should always use BOTH data structures. Both are needed for the correct executing of the program.

Recall from earlier:




Figure 4: View of Runtime Stack as code is executed.

Initially the `framePointer` stack will have 0 for step 1 above. When we call function `f` at step 2, `framePointer` stack will have the values 0 and 3 since 0 is the start of `main` and 3 is the start of `f`. At step 4.1 we pop the `framePointer` stack (3 is popped). Then the `framePointer` Stack

will only contain the value 0. This is Mains starting frame index. Note when transitioning from step 4 and 4.1 there is going to be more work that is needed to be done to clean up each frame before returning from functions.

7.1 RunTimeStack Class Functions

```
1 /**
2  * Used for dumping the current state of the runTimeStack.
3  * It will print portions of the stack based on respective
4  * frame markers.
5  * Example [1,2,3] [4,5,6] [7,8]
6  * Frame pointers would be 0,3,6
7  */
8 public String dump() { }
9
10 /**
11  * returns the top of the runtime stack, but does not remove
12  * @return copy of the top of the stack.
13  */
14 public int peek(){ }
15
16 /**
17  * push the value i to the top of the stack.
18  * @param i value to be pushed.
19  * @return value pushed
20  */
21 public int push(int i) { }
22
23 /**
24  * removes to the top of the runtime stack.
25  * @return the value popped.
26  */
27 public int pop() { }
28 /**
29  * Takes the top item of the run time stack, and stores
30  * it into a offset starting from the current frame.
31  * @param offset number of slots above current frame marker
32  * @return the item just stored
33  */
34 public int store(int offsetFromFramePointer) { }
35 /**
36  * Takes a value from the run time stack that is at offset
37  * from the current frame marker and pushes it onto the top of
```

```
38  * the stack.
39  * @param offset number of slots above current frame marker
40  * @return item just loaded into the offset
41  */
42 public int load(int offsetFromFramePointer) { }
43
44 /**
45  * create a new frame pointer at the index offset slots down
46  * from the top of the runtime stack.
47  * @param offset slots down from the top of the runtime stack
48  */
49 public void newFrameAt(int offsetFromTopOfRunStack) { }
50
51 /**
52  * pop the current frame off the runtime stack. Also removes
53  * the frame pointer value from the FramePointer Stack.
54  */
55 public void popFrame() { }
```

8 VirtualMachine Class

VirtualMachine Class is used for executing the given program. The VirtualMachine basically acts as the controller of this program. All operations need to go through the VirtualMachine class. It will contain the following data-fields (more can be added IF needed):

```
1 private RunTimeStack    runTimeStack;
2 private Stack<Integer> returnAddress;
3 private Program         program;
4 private int              programCounter;
5 private boolean          isRunning;
```

The VirtualMachine will contain many functions. These will not be listed. The reason for this is that you are expected to abstract the operations needed by the ByteCodes and create clean and concise set of functions. Do your best to limit the amount of duplicate code. Points will be taken away from VirtualMachines that contain a lot of duplicate code (or simply creating a function for each byte code). Also, limit the amount of ByteCode logic that appears in the VirtualMachine. For example, if you are executing the ReadCode ByteCode, the process of reading is not done by the VirtualMachine, it is done by the ReadCode class and then the ReadCode class requests the VirtualMachine to push the read value onto the RuntimeStack.

The returnAddress stack stores the ByteCode index(PC) that the VirtualMachine should execute when the current function exits. Each time a function is entered, the PC should be pushed onto the returnAddress stack. When a function exits the PC should be restored to the value that is popped from the top of the returnAddress Stack.

One important function in the VirtualMachine is execute program. A sample base function is given:

```
1 public void executeProgram(){
2     isRunning = true;
3
4     while(isRunning){
5         ByteCode code = program.getCode(programCounter);
6         code.execute(this);
7         programCounter++;
8     }
9 }
```

Note that we can easily add new ByteCodes without breaking the operation of the executeProgram function. We are using Dynamic Binding to achieve code flexibility, extensibility and readability. This loop should be very easy to read and understand.

9 Interpreter Dumping

There is one special ByteCode that is used to determine whether the VirtualMachine should dump the current state of the RunTimeStack and the information about the currently executed ByteCode. This ByteCode is named DUMP. It will have 2 states, ON and OFF. No other form of DUMP will occur in this program or any given program. Note that your test files(.x.cod) may not contain these and you will need to add them yourself.

- DUMP ON is an interpreter command to turn on runtime dumping. This will set an interpreter switch that will cause runTimeStack dumping AFTER execution of EACH ByteCode. This switch or variable will be stored as a private data-field in the VirtualMachine class.
- DUMP OFF will reset the switch to end dumping. Note that DUMP instructions will not be printed.
- DO NOT dump program state unless dumping is turned ON.

Consider the following ByteCode program (ByteCodes marked with arrows are dumped):

```
GOTO start<<1>>
LABEL Read
REUTRN
LABEL Write
LOAD 0 dummyformal
WRITE
RETURN
LABEL start<<1>>
LIT 0 i
LIT 0 j
GOTO continue<<3>>
LABEL f<<2>>          <-- Dumped
LIT 0 j              <-- Dumped
LIT 0 k              <-- Dumped
LOAD 0 i             <-- Dumped
LOAD 0 j             <-- Dumped
DUMP OFF             DUMP IS TURNED OFF
BOP +
LOAD 2 k
BOP +
LIT 2
BOP +
RETURN f<<2>>
POP 2
```

LIT 0 GRATIS-RETURN-VALUE	
RETURN f<<2>>	
LABEL continue<<3>>	
DUMP ON	DUMP IS TURNED ON
LIT 0 m	<-- Dumped
LIT 3	<-- Dumped
ARGS 1	<-- Dumped
CALL f<<2>>	<-- Dumped
DUMP ON	DUMP IS TURNED ON
STORE 2	<-- Dumped
DUMP OFF	DUMP IS TURNED OFF
LOAD 1 j	
LOAD 2 m	
BOP +	
ARGS 1	
WRITE	
STORE 0 i	
POP 3	
HALT	

When dumping is turned on you should print the following information **JUST AFTER** executing the current ByteCode:

- Print the ByteCode that was just executed (DO NOT PRINT the DUMP ByteCodes)
- Print the runtime stack with spaces separating frames (just after the ByteCode was executed). Values contained in one frame should be surrounded by [and].
- If dump is not on then DO NOT print the ByteCode, nor dump the runtime stack.

```

LIT 0 m      int m
      [0,0,0]
LIT 3
      [0,0,0,3]
ARGS 1
      [0,0,0] [3]
CALL f<<2>>  f(3)
      [0,0,0] [3]
LIT 0 j      int j
      [0,0,0] [3,0]
LIT 0 k      int k

```

```

        [0,0,0] [3,0,0]
LOAD 0 i <load i>
        [0,0,0] [3,0,0,3]
LOAD 1 j <load j>
        [0,0,0] [3,0,0,3,0]
...
STORE 2 m      m = 2
        [0,0,5]

```

Note : Following shows the output if dump is on and 0 is at the top of the runtime stack. RETURN f<< 2 >> exit f:0 note that 0 is returned from f<< 2 >>

- If Dumping is turned on and we encounter an instruction such as WRITE then output as usual; the value may be printed either before or after dumping information e.g.,

```

LIT 3
    [0,0,0,3]
3
WRITE
    [0,0,0,3]

```

9.1 DUMPING IMPLEMENTATION NOTES

The Virtual Machine maintains the state of your running program, so it is a good place to have the dump flag. You should not use a state variable in the ByteCode class.

The dump method should be a part of the RunTimeStack class. This method is called without any arguments. Therefore, there is no way to pass any information about the VirtualMachine or the ByteCode classes into RunTimeStack. As a result, you can't really do much dumping inside RunTimeStack.dump() except for dumping the state of the RunTimeStack itself. Also, **The ByteCode abstraction and its concrete classes should not be calling dump in the VM. Nor should they be asking the VirtualMachine if dumping is on. The VirtualMachine knows all the information to trigger dumping. Each ByteCode SHOULD implement a method to support dumping of the bytecode.**

It is impossible to determine the declared type of a variable by looking at the ByteCode file. To simplify matters you should assume whenever the interpreter encounters LIT 0 x (for example), that it is an int. When you are dumping the ByteCodes, it is ok to represent the values as int.

10 Coding Hints

10.1 Possible Order of Implementation

Below is an ordered list of how you may start this assignment. Note that this is not the only way. If you have a better way you may use your way.

1. **Read this PDF 3 or 4 times minimum.**
2. Think and design the ByteCode abstraction. This will help with completing other classes.
3. Implement and test the RuntimeStackStack class.
4. Create all ByteCode classes listed in the table of ByteCodes listed above. The classes do not need to be fully implemented but having the method stubs (empty method bodies) is enough.
5. Implement the ByteCodeLoader class. Note as you implement the loadCodes function, you will need to start filling in the init functions for each of the ByteCode classes as well as adding helper functions to the Program class.
6. Implement the Program Class, specifically the resolveAddress function.
7. Implement the VirtualMachine Class.
8. While completing the implementation of the VirtualMachine class, you should also start to fill in the execute method of each of the ByteCode classes.

10.2 Other Coding Hints

Below are general coding hints and assumptions you may make about the project and the source files used.

- You will assume that the given ByteCode source programs (.cod files) you will use for testing are generated correctly and contain NO ERRORS.
- You will need to make sure that you protect the RunTimeStack from stack overflow or stack underflow errors. Also make sure no ByteCode can pop past any frame boundary.
- You should provide as much documentation as you can when implementing The Interpreter. Short, OBVIOUS functions don't need comments. However, you should comment each class describing its function and purpose. Take in consideration that your code is the first level of documentation. How you name things in your code matters. If you find yourself writing a lot of comments, then you may be either explaining a complex algorithm, which is OK, or your code contains named variable and methods that are not descriptive enough.

- **DO NOT** provide any methods that return any components contained WITHIN the VirtualMachine(this is the exact situation that will break encapsulation) – you should request the VirtualMachine to perform operations on its components. This implies that the VirtualMachine owns the components and is free to change them as needed without breaking clients' code (for example, suppose I decided to change the name of the variables that holds my runtime stack - if your code had a reference to that variable then your code would break. This is not an unusual situation – you can consider the names of methods in the Java libraries that have been marked deprecated).

The only downside is it might be a bit inefficient. Since I want to impress on everyone important software engineering issues, such as encapsulation benefits, I want to enforce the requirement that you **DO NOT BREAK** encapsulation.

Consider that the VirtualMachine calls the individual ByteCodes' execute method and passes itself as a parameter. For the ByteCode to execute, it must invoke 1 or more methods in the runStack object. It can do this by executing VirtualMachine.runStack.pop(); however, this DOES break encapsulation. To avoid this, you will need to have a corresponding set of methods within the VirtualMachine that do nothing more than pass the call to the runStack. For example, a pop function can be implemented in the VirtualMachine in the following way

```
1 public int popRunStack() {  
2     return runStack.pop();  
3 }
```

Then in a ByteCodes' execute method a pop can be executed as:

```
1 int temp = VirtualMachine.popRunStack();
```

- Each ByteCode class should have fields for their specific arguments. The ByteCode abstraction and concrete classes **SHOULD NOT CONTAIN ANY FIELDS (instance variables) THAT ARE RELATED TO STORING ARGUMENTS AS A COLLECTION OR ARRAY. CONCRETE BYTECODES MUST STORE ARGUMENTS AS DATA-FIELDS AND THEY MUST THE CORRECT TYPE.** This is a design requirement.

It is easier to think in more general terms (i.e. plan for any number of arguments for a ByteCode). Note that the ByteCode abstraction should not be aware of the peculiarities of any ByteCode. That is, some ByteCodes might have zero arguments (HALT) or one argument, etc. Consider providing an init function with each ByteCode class. After constructing a ByteCode instance during the loadCodes phase, you can then call init passing in an ArrayList of String as arguments. Each ByteCode object will then interrogate the ArrayList and extract the needed arguments itself. The ByteCode abstraction should not record any arguments for any ByteCodes. Each ByteCode

concrete class will need instance variables for each of their arguments. There may be a need for additional instance variables as well.

- When you read a line from the ByteCode file, you should parse the arguments and place them in an ArrayList of Strings. Then pass this ArrayList to the ByteCode's init function. Each ByteCode is responsible for extracting relevant information from the ArrayList and storing it as private data.
- **The ByteCode abstraction and its concrete classes must be contained in a package named ByteCode.**
- Any output produced by the WRITE ByteCode will be interspersed with the output from dumping (if dumping is turned on). In the WRITE action you should print one number per line. DO NOT print out something like :

Program Output: 2

Instead, only need to include the value and that is it. It should just be:

2

- There is no need to check for division by zero error. Assume that you not have a test case where this occurs.

11 Setup

11.1 Importing Project

When importing The interpreter project you will use the root of your repository as the source. **DO NOT USE** the interpreter folder as the root. This will cause the project to not work and in the end force you to change the file structure. As has been noted many times, **YOU ARE NOT ALLOWED TO MODIFY THE THE BASE STRUCUTRE OF THIS PROJECT. DOING WILL CUASE A POINT PENLTY TO BE APPLIED TO YOUR GRADE.**

When executing this project you will need to create the run configurations and set the command line arguments. The Interpreter only is able to handle the .x.cod files. The .x are not going to be used for this project, and are present only as a reference.

A sample picture of a run configuration is given below:

Note that the working directory and the Use Classpath of module values may not be the same. The purpose of the image it show where the filename of the ByteCode source file goes.




Figure 5: Sample picture of a IntelliJ run configuration for The Interpreter project.

12 Requirements

1. Implement ALL supported bytecodes classes listed in the [Supported ByteCodes](#). Be sure to create the correct abstractions for the bytecodes. It is possible to have multiple abstractions within the set of byte codes classes.
2. Make sure to adhere to the design requirements listed in the above sections for each of the classes and the ByteCode abstraction.
3. Complete the implementation of the following classes:
 - a ByteClassLoader
 - b Program
 - c RuntimeStack
 - d VirtualMachine

The Interpreter class have already been implemented for you. Interpreter.java is the entry point to this project. All projects will be graded using this entry point. **The Interpreter class MUST NOT be changed. Doing so will cause you project to lose points.**

4. Make sure that all variables have their correct modifiers. Projects with variables using the incorrect modifiers will lose points.

5. Make sure not to break encapsulation. Projects that contain objects or classes trying to access members that it should not be allowed to will lose points. For example, if a ByteCode needs to access or write to the runtime stack, it should **NOT** be allowed to. It needs to request these operations from the virtual machine. Then the virtual machine will carry out the operation.

It would also be incorrect to make a method in the virtual machine for each byte code. While this approach will work, this solution will produce a lot of duplicate code. Points will be deducted for this type of solution. Do your best to understand the operations of each ByteCode and how they manipulate the data-structures the Virtual Machine maintains. You will be able to see that some ByteCodes operate on these data-structures in a very similar way. Basically, you are being asked to code to the virtual machine and not to the byte codes.

6. The basic structure of the interpreter folder and its contents cannot be changed. You may add sub-folders to the ByteCode folder. **10 POINTS WILL BE DEDUCTED if the structure of the project is changed.**

13 Submission

1. All completed source code needs to be submitted to your repository created with the link on iLearn. Use the following commands to commit your code. Please note you must be in the repository folder to run these commands.

```
1 $ git add .
2 $ git commit -m "message"
3 $ git push origin master
```

2. Complete the required documentation as outlined in the documentation guidelines. This is stored in the documentation folder. Convert the completed document to a PDF file and save in the documentation folder. **PDF FILES ARE REQUIRED. 5 POINTS WILL BE DEDUCTED FOR NOT SUBMITTING A PDF**